



UNIVERSITY OF
ILLINOIS
URBANA - CHAMPAIGN

Batch Inference Load Balancing for Mixture-of-Experts LLMs*

Peter Lin

*Idea originally from group research project for CS 498 ML Systems

2/16/2026

Project Overview: Optimizing Machine Learning Systems



Background and Motivation: What is MoE? Why use batch inference?

Problem: How does batching impact expert load balance?

Proposed Solution: How does the proposed algorithm address the core of the problem?

Experimental Setup: Hardware details and evaluation dataset

Initial Evaluation: What are the baseline statistics?

Final Evaluation: How does the optimized model perform compared to baseline?

Conclusion: Final thoughts and takeaways

Background: Mixture-of-Experts (MoE) Models are Sparse



In dense models, all parameters are activated during generation

- Computationally expensive

In sparse models, only a subset of parameters are activated

- Computationally efficient
- Uses a router to assign specific tokens to a subset of (k) experts
 - Top- k experts chosen based on score




Image: <https://www.maartengrootendorst.com/blog/moe/>

Background: Qwen3 MoE Architecture



Qwen3-30B-A3B has decoder-only architecture

- 48 decoder layers
- 128 unique experts per layer
- $k = 8$ for top-k selection

We can interact with intermediate computations

- e.g. hook into gate module to see router's top-k expert selection
- e.g. replace the forward function in the mlp module with a custom forward function

```
Qwen3MoeForCausalLM(  
    (model): Qwen3MoeModel(  
        (embed_tokens): Embedding(151936, 2048)  
        (layers): ModuleList(  
            (0-47): 48 x Qwen3MoeDecoderLayer(  
                (self_attn): Qwen3MoeAttention(  
                    (q_proj): Linear(in_features=2048, out_features=4096, bias=False)  
                    (k_proj): Linear(in_features=2048, out_features=512, bias=False)  
                    (v_proj): Linear(in_features=2048, out_features=512, bias=False)  
                    (o_proj): Linear(in_features=4096, out_features=2048, bias=False)  
                    (q_norm): Qwen3MoeRMSNorm((128,), eps=1e-06)  
                    (k_norm): Qwen3MoeRMSNorm((128,), eps=1e-06)  
                )  
                (mlp): Qwen3MoeSparseMoeBlock(  
                    (experts): Qwen3MoeExperts(  
                        (act_fn): SiLUActivation()  
                    )  
                    (gate): Qwen3MoeTopKRouter()  
                )  
                (input_layernorm): Qwen3MoeRMSNorm((2048,), eps=1e-06)  
                (post_attention_layernorm): Qwen3MoeRMSNorm((2048,), eps=1e-06)  
            )  
            (norm): Qwen3MoeRMSNorm((2048,), eps=1e-06)  
            (rotary_emb): Qwen3MoeRotaryEmbedding()  
        )  
        (lm_head): Linear(in_features=2048, out_features=151936, bias=False)  
    )
```

All 128 experts live here

Router lives here

Motivation: Batching Inputs Leverages GPU Parallelism



Without batching, individual inputs are processed sequentially



Offline batching inputs enables parallel processing but introduces pad tokens (**focus of this project**)



Continuous batching packs inputs tighter but introduces infra overhead (e.g. vLLM)



Image: <https://www.baseten.co/blog/continuous-vs-dynamic-batching-for-ai-inference/>

Problem: Offline Batch Inference Induces Expert Load Imbalance



Heatmaps plot token distribution across every expert in the model

- All 128 experts in all 48 layers (6,144 unique experts)


No batching

- Tokens are relatively evenly distributed


With batching

- Pad tokens clearly overpower the distribution

Token Distribution Heatmap (No Batching)



Token Distribution Heatmap (With Batching, Same Data)




Solution: Reroute Padding Tokens to Underutilized Experts




Optimized forward function:

1. Get padding mask (1 if padding 0 otherwise)
2. Get expert selection from gate network
3. Compute padding indices and number of pad tokens
4. If number of pad tokens is positive:
 - a. Redistribute every pad token's expert assignments to bring distribution closer to average “real” token count
 - b. Update expert selection
5. Compute expert outputs using expert selection
6. Return final tensor

e.g. rerouting pad tokens from experts 1 and 4 across all 6 experts



Observation: Sorting Before Batching Reduces Padding



Experimental Setup



Hardware: 1xA100 GPU (Google Colab, High-RAM)

Model: Qwen3-30B-A3B

Dataset: Massive Multitask Language Understanding (MMLU)

Baseline evaluation on sorted and unsorted batches

Final evaluation comparing baseline with optimization


For each evaluation:

- Measure load balance with coefficient of variation (CV), defined as std dev / mean
- Measure execution time based on time taken to complete processing a single batch
- Measure model output quality using MMLU accuracy


Measure stats for four runs (batch size x batch count):

- 8x50, 16x20, 32x10, 64x5

Initial Evaluation: Baseline (Unsorted and Sorted Batches)



Final Evaluation: Optimization vs Baseline





Conclusion: Rerouting Padding Tokens Improves Load Balance



As batch size increases, coefficient of variation gets lower when padding tokens are rerouted

The load-balance-optimized forward function is efficient; execution time only slightly increases and accuracy is comparable to baseline

The next step is to test the algorithm on a multi-GPU node and measure latency improvements



Questions?

